

# When Does Post-Hoc $K=V$ Tying Fail?

## A Convex-Hull Constraint in Self-Attention

Companion note to *Do Transformers Need Three Projections?* (Kayyam et al.)

### Abstract

*Do Transformers Need Three Projections?* shows that a  $Q \neq K=V$  attention variant, trained *from scratch*, matches full QKV attention on a range of synthetic and vision tasks. This note asks a different question: what happens if  $K$  and  $V$  are tied *after* training, by surgically merging an already-trained QKV model’s independent key and value projections, with no retraining? The central claim is narrow and architecture-level, not a claim about overall transformer expressivity: once  $K = V$ , a *single attention sub-block’s* output is confined to the convex hull of the shared key/value vectors, regardless of the attention weights (Proposition 1); the residual stream, MLPs, and later layers are unrestricted by this and can and do carry information outside that hull. This predicts, and we confirm across five synthetic sequence tasks and four vision-classification datasets, a sharp split: the two synthetic tasks solvable without cross-token content routing are completely unaffected by the surgery, while every task or dataset that requires moving one position’s value to another collapses toward chance. We then test the natural mechanistic story for the surgery-dependent cases – that the teacher’s  $K$ -projection specialized for *addressing* at the expense of *content*, while  $V$  specialized the other way – and find it is only half right: a linear probe recovers token identity from  $K$  and  $V$  activations equally well, so the failure is not a loss of content information in  $K$ . A systematic interpolation between  $K$  and  $V$ , together with a random-projection control, shows that reusing  $K$  (KEEP\_K) is often no better than an unrelated, untrained projection despite retaining linearly decodable token information – pointing to a representation-compatibility account rather than a simple information-loss one, though we have not measured downstream calibration to  $K$  vs.  $V$  directly and hold this as the best-supported reading rather than an established mechanism. Collapse severity also correlates with how functionally redundant the trained  $K/V$  projections already were (linear CKA on held-out activations,  $r = -0.84$ ,  $n = 9$ , though this drops to  $\rho = -0.56$  (n.s.) under Spearman and is not separately significant within either task-type group at  $n = 5/n = 4$ , so we treat CKA as a useful diagnostic rather than a validated general predictor) – and we surface, without resolving, a puzzle inside even that: within vision,  $K$ - $V$  similarity *rises* rather than falls as classification difficulty increases from MNIST to CIFAR-100, even as collapse severity rises too. Finally, the collapse is comparatively cheap to repair: knowledge distillation from the original QKV model recovers the from-scratch tied baseline in 14 of 15 synthetic (task, surgery-mode) configurations within 5 epochs, and on vision, AVG-mode surgery reaches (and on CIFAR-10 slightly exceeds) the from-scratch tied baseline within 10 epochs, while KEEP\_K remains the one mode with a persistent, slow-converging gap – consistent with the basis-mismatch account, since it is the mode with the most re-calibration to do. Our overarching point is not that  $K = V$  is fundamentally deficient: jointly training  $Q \neq K=V$  from scratch works fine (that is the base paper’s headline result), and untying is trivial for the two tasks that never needed cross-token routing. The point is that *post-hoc merging of independently-trained representations is a different operation from a joint architectural constraint*, with its own, separately characterizable failure mode.

# 1 Setup and notation

**Relationship to the base paper.** *Do Transformers Need Three Projections?* asks whether  $Q \neq K = V$ , trained jointly from a random initialization, can match full QKV attention – and finds that it can. This note asks a different question about the same equality: can the representations learned by *independent*  $K$  and  $V$  projections, already specialized apart over a full training run, be merged *after the fact*? These are not the same intervention. Joint training lets addressing and content pressures negotiate a single shared representation from the start, under gradient descent’s full control; post-hoc merging takes two representations optimized under no such constraint and imposes it externally, with no opportunity for the rest of the network to adapt. The central scientific question here is specifically the latter: whether, and how predictably, independently-trained representations resist being merged – not whether  $K = V$  is expressive enough when it is a design choice made before training starts.

We use the attention parameterization of the base paper: for input  $x \in \mathbb{R}^{T \times d}$ , single-head<sup>1</sup> self-attention computes

$$q_i = c_q(x_i), \quad k_j = c_k(x_j), \quad v_j = c_v(x_j), \quad a_{ij} = \text{softmax}_j\left(\frac{q_i \cdot k_j}{\sqrt{d}}\right), \quad y_i = \sum_{j=1}^T a_{ij} v_j,$$

where  $c_q, c_k, c_v$  are learned linear maps. The base paper’s six variants are recovered by sharing these maps: QKV (none shared),  $Q=K \neq V$  ( $c_q = c_k$ ),  $Q \neq K=V$  ( $c_k = c_v$ ), and  $Q=K=V$  (all three shared), each optionally with a 2D positional injection (the “+” variants). Every block also carries a residual connection,  $h_i^{\text{out}} = h_i^{\text{in}} + \text{Attn}(\cdot)_i$ , followed by a position-wise MLP with its own residual connection.

**Surgery.** Given a *trained* QKV checkpoint (independent  $c_k, c_v$ ), we construct a  $Q \neq K=V$  student with no retraining by setting the student’s shared projection  $c_{kv}$  to one of:

- KEEP\_K:  $c_{kv} \leftarrow c_k$  (drop the trained  $V$  entirely),
- KEEP\_V:  $c_{kv} \leftarrow c_v$  (drop the trained  $K$  entirely),
- AVG:  $c_{kv} \leftarrow \frac{1}{2}(c_k + c_v)$  (average the two Linear layers’ weights and biases; since both are affine maps this is identical to averaging their outputs).

All other parameters ( $c_q$ ,  $c_{\text{proj}}$ , MLPs, layer norms, embeddings, head) are copied verbatim from the teacher.

## 2 A convex-hull obstruction

**Proposition 1** (Convex-hull confinement under  $K=V$ ). *If  $k_j = v_j$  for every position  $j$  (as holds for all three surgery modes above, and for the from-scratch  $Q=K=V$  and  $Q \neq K=V$  variants), then for every query position  $i$ ,*

$$y_i \in \text{conv}\{k_1, \dots, k_T\} \subseteq \text{span}\{k_1, \dots, k_T\},$$

*regardless of the attention weights  $a_{ij}$ .*

*Proof.*  $a_{ij} \geq 0$  and  $\sum_j a_{ij} = 1$  by construction of the softmax, so  $y_i = \sum_j a_{ij} k_j$  is by definition a convex combination of  $\{k_j\}$ .  $\square$

---

<sup>1</sup>The multi-head case is identical after splitting  $d$  into  $H$  blocks of size  $d/H$ ; the argument in Section 2 applies unchanged within each head.




Figure 1: The same architecture,  $K=V$ , reached two different ways. Left: the QKV teacher, with independent  $K/V$  projections free to specialize for addressing and content separately. Middle: post-hoc surgery forces one already-specialized representation to do both jobs, with nothing else in the network adjusted (Section 4.2 finds the result closer to a basis mismatch than to information loss). Right: choosing  $K=V$  before training lets every other parameter co-adapt to it from the start – the same mathematical bound (Proposition 1) applies, but nothing has to unlearn a different assumption.

**Corollary 1.** *The attention sub-block can only make available to position  $i$  information that is already representable inside  $\text{span}\{k_1, \dots, k_T\}$ . If the trained  $K$ -projection was shaped by gradient descent to produce a good addressing signal ( $q_i \cdot k_j$  ranking the correct target highest) without pressure to remain injective in the token content  $x_j$  — because that pressure was carried entirely by the independent  $V$ -projection in the QKV teacher — then  $k_j$  need not determine  $x_j$ , and routing  $k_j$  itself to position  $i$  can lose the very content the downstream task needs. This is a hypothesis about why surgery hurts, stated here before we test it; Section 4.2 probes addressing and content recoverability directly and finds the picture is more specific than “ $K$  lacks content” as stated — content is present in  $K$  too, but evidently in a form the rest of the network cannot use.*

**Remark 1** (Invariance to normalization). *Proposition 1’s proof uses only that  $a_{ij} \geq 0$  and  $\sum_j a_{ij} = 1$ ; it places no assumption on how  $k_j$  (equivalently  $v_j$ , since  $k_j = v_j$ ) is computed. It therefore holds unchanged whether  $k_j$  is a raw linear projection, an  $\ell_2$ -normalized projection (as in QK-normalized attention variants), or the output of any other fixed transformation applied before the weighted sum — confinement is to the convex hull of whatever vectors are actually summed, in whatever geometry they live in. In the checkpoints studied here, LayerNorm is applied pre-block to the residual stream  $x_j$  before  $c_k, c_v$  project it (Section 1), not to  $k_j$  or  $v_j$  themselves, so this is a non-issue for our experiments specifically; we state the general invariance here because it is easy to construct architectures where it would otherwise look like a threat to the argument. A normalization that discarded information  $v_j$  needed to retain would show up as a smaller effective  $\text{span}\{k_j\}$  in Corollary 1’s premise, not as a failure of Proposition 1 itself, which is unconditional.*

**Remark 2** (Scope: one attention sub-block, not the whole transformer). *Proposition 1 constrains  $y_i$ , the output of one attention sub-block. It does not constrain  $h_i^{\text{out}} = h_i^{\text{in}} + y_i$ , the MLP applied to it, or any later layer: the residual stream carries  $h_i^{\text{in}}$  forward regardless of what the attention sub-block computes, and nothing here restricts what a downstream MLP or a later attention layer (with its own, unconstrained  $K'/V'$ ) can do with the result. A model with  $K = V$  in every layer is therefore not restricted to representing only what lies in a single layer’s convex hull — claiming*

that would be a claim about the whole network’s expressivity, which Proposition 1 does not make and we do not intend. What the proposition does say is narrower and, we think, still useful: a single  $K = V$  attention operation cannot independently choose the representation used for content transport from the representation used for addressing – what gets transported to position  $i$  must be a convex combination of the very vectors  $\{k_j\}$  that determined where to look, not some separately-chosen payload. Whether that narrower constraint ends up mattering for a given task is exactly the empirical question the rest of the note is about.

Proposition 1 is therefore silent on *when* its obstruction actually bites in practice: a task that never needs cross-token content routing in the first place is immune to it regardless of how many layers have  $K = V$ . Section 3 makes this precise for our tasks; Section 4 shows that even among the routing-dependent tasks, the practical severity varies with how much the teacher’s  $K$  and  $V$  projections had actually specialized (Corollary 1’s premise) – which we measure and test directly rather than assume, and which turns out to need a more specific account than the corollary as originally stated (Section 4.2).

### 3 Why pointwise tasks are immune

Two of the five synthetic tasks are pointwise: SUB computes  $y_i = 9 - x_i$  and COPY computes  $y_i = x_i$ , each a function of position  $i$ ’s own input only. Each block’s residual connection preserves a direct, additive path from position  $i$ ’s own incoming representation ( $\text{embed}(x_i) + \text{pos\_emb}(i)$  at the first block, that block’s own input thereafter) through to its output, regardless of what the attention sub-block computes; stacked over every block, that representation is never overwritten, only added to. Combined with per-position MLPs that are already universal approximators of pointwise maps, attention’s specific content-routing is therefore not load-bearing for these two tasks, and Proposition 1’s obstruction is vacuous: it restricts what attention can contribute, but the task does not need attention to contribute anything.

This is not merely a plausibility argument – it is directly testable with data we already have. If attention’s specific query–key matching mattered for SUB/ COPY, corrupting the value stream with either the wrong projection (KEEP\_K) or corrupting the key stream with the wrong projection (KEEP\_V) should hurt *at least one* of the two directions. Instead, Table 1 shows *both* surgeries leave SUB and COPY at 1.000 accuracy, exactly as the residual-bypass account predicts, while the three routing-dependent tasks (REVERSE, SORT, SWAP – each requires moving a *different* position’s value to position  $i$ ) collapse sharply under the same surgeries.

## 4 Empirical validation

### 4.1 Zero-shot surgery and distillation recovery

Table 1 and Table 2 report, for each task/dataset: the from-scratch QKV teacher accuracy, the from-scratch tied ( $Q \neq K = V$ ) baseline accuracy (a reference performance a distilled student can reasonably target – not a hard ceiling, since AVG/CIFAR-10 slightly exceeds it (see Table 2)), the zero-shot accuracy immediately after surgery (no fine-tuning), and the accuracy after 5 epochs of distillation from the QKV teacher (cross-entropy on hard labels + KL divergence against the teacher’s soft logits,  $T = 2$ , equal weighting).

AVG is uniformly the safer zero-shot surgery on every vision dataset, and doubling the distillation budget does not treat all three surgery modes equally: KEEP\_K is the one mode with a persistent

| Task    | Mode   | QKV   | Q $\neq$ K=V (scratch) | Zero-shot surgery | Distilled (5 ep) |
|---------|--------|-------|------------------------|-------------------|------------------|
| REVERSE | keep_k | 1.000 | 1.000                  | 0.100             | 1.000            |
|         | keep_v | 1.000 | 1.000                  | 0.123             | 1.000            |
|         | avg    | 1.000 | 1.000                  | 0.861             | 1.000            |
| SORT    | keep_k | 0.998 | 0.996                  | 0.583             | 0.965            |
|         | keep_v | 0.998 | 0.996                  | 0.717             | 0.996            |
|         | avg    | 0.998 | 0.996                  | 0.751             | 0.996            |
| SUB     | any    | 1.000 | 1.000                  | 1.000             | 1.000            |
| SWAP    | keep_k | 1.000 | 1.000                  | 0.088             | 1.000            |
|         | keep_v | 1.000 | 1.000                  | 0.122             | 1.000            |
|         | avg    | 1.000 | 1.000                  | 0.906             | 1.000            |
| COPY    | any    | 1.000 | 1.000                  | 1.000             | 1.000            |

Table 1: Synthetic sequence tasks (token accuracy). Routing-dependent tasks (REVERSE, SORT, SWAP) collapse under surgery and recover under distillation; pointwise tasks (SUB, COPY) are unaffected by either. 14 of these 15 (task, surgery-mode) configurations reach the from-scratch tied baseline within 5 epochs; the one exception, SORT/KEEP\_K, reaches 0.965 against a 0.996 target and is still improving at epoch 5.

gap even at 10 epochs (Section 4.2 tests why, on the synthetic tasks, and finds the evidence more consistent with a basis-mismatch account than a simple content-loss one).

Tables 1–2 report only endpoint accuracies; Figure 2 shows the full per-epoch recovery curve for every (task or dataset, surgery-mode) configuration referenced above. The shape confirms what the endpoints suggest: REVERSE/SUB/SWAP/COPY and MNIST are already flat by epoch 1 (the surgery was mild or the gap closes almost immediately), SORT/KEEP\_K and every CIFAR/KEEP\_K curve are still rising through their full epoch budget without having visibly plateaued as early as epoch 2–3, and KEEP\_V/AVG on CIFAR flatten out well before their budget ends – consistent with KEEP\_K being a harder optimization problem rather than merely a deeper hole to climb out of.

**Where does the zero-shot optimum sit, and is reusing  $K$  better than noise?** Tables 1–2 report only three points along the line  $c_{kv} = \alpha W_K + (1 - \alpha)W_V$  ( $\alpha=1$ : KEEP\_K,  $\alpha=0$ : KEEP\_V,  $\alpha=0.5$ : AVG). Figure 3 sweeps  $\alpha \in \{0, 0.1, \dots, 1.0\}$  and adds a fourth reference point:  $c_{kv}$  left at its fresh, untrained initialization, with no relationship to the teacher’s  $K$  or  $V$  at all – reported as a mean  $\pm 1$  std. over 5 random seeds, not a single number, since the baseline’s own sampling variance turns out to matter for what we can claim about it.

Two things stand out. First, the curve is neither monotonic nor symmetric around AVG: REVERSE, SWAP, and all four vision datasets peak near  $\alpha \approx 0.4$ – $0.5$ , while SORT peaks earlier, around  $\alpha \approx 0.2$ – $0.3$ , and degrades monotonically from there to KEEP\_K. SUB and COPY are flat at 1.000 across every  $\alpha$ , including every random seed – a stronger confirmation of Section 3 than the three original discrete points.

Second: KEEP\_K does not reliably preserve useful downstream functionality despite preserving linearly decodable content (Section 4.2). Table 3 compares KEEP\_K against the random baseline’s mean and std. for each non-pointwise task. On SWAP, KEEP\_K sits about 12 random- baseline standard deviations below the mean – unambiguous. On the other five (REVERSE, SORT, MNIST, FMNIST, CIFAR-10, CIFAR-100), KEEP\_K sits within about 2 standard deviations of the random baseline in either direction, which at  $n = 5$  seeds we cannot call a reliable difference from noise in either direction. We had originally read the point estimates alone as “KEEP\_K is usually at or below random,” which overstated what 5 seeds actually support; the defensible claim is narrower and, we think, still the interesting one: outside of SWAP’s dramatic case, KEEP\_K’s specialized

| Dataset   | Mode   | QKV   | Q $\neq$ K=V (scratch) | Zero-shot surgery | Distilled (5 ep) | Distilled (10 ep) |
|-----------|--------|-------|------------------------|-------------------|------------------|-------------------|
| MNIST     | keep_k | 0.981 | 0.978                  | 0.114             | 0.979            | –                 |
|           | keep_v | 0.981 | 0.978                  | 0.151             | 0.982            | –                 |
|           | avg    | 0.981 | 0.978                  | 0.910             | 0.983            | –                 |
| FMNIST    | keep_k | 0.887 | 0.882                  | 0.084             | 0.881            | –                 |
|           | keep_v | 0.887 | 0.882                  | 0.105             | 0.883            | –                 |
|           | avg    | 0.887 | 0.882                  | 0.811             | 0.888            | –                 |
| CIFAR-10  | keep_k | 0.696 | 0.698                  | 0.076             | 0.631            | 0.666             |
|           | keep_v | 0.696 | 0.698                  | 0.153             | 0.686            | 0.693             |
|           | avg    | 0.696 | 0.698                  | 0.368             | 0.696            | <b>0.702</b>      |
| CIFAR-100 | keep_k | 0.437 | 0.445                  | 0.007             | 0.384            | 0.411             |
|           | keep_v | 0.437 | 0.445                  | 0.021             | 0.424            | 0.430             |
|           | avg    | 0.437 | 0.445                  | 0.109             | 0.432            | 0.435             |

Table 2: Vision classification (ViT, top-1 accuracy). Unlike the synthetic tasks, every dataset shows some collapse under KEEP\_K/KEEP\_V: these ViTs classify from a learned CLS token ( $\mathbf{h}[:, 0]$  after the final block), not global average pooling, so which patches reach the classification decision is itself an attention-weighted, content-routing operation rather than a fixed uniform average – there is no pointwise bypass analogous to Section 3. At the 5-epoch budget, KEEP\_K on CIFAR-10 and CIFAR-100 was still visibly improving, so both CIFAR datasets (all three modes) were rerun with a 10-epoch budget; MNIST/FMNIST were not rerun since they were already converged by 5 epochs. KEEP\_V and AVG are essentially converged on CIFAR by 10 epochs (AVG/CIFAR-10 slightly exceeds the from-scratch tied baseline); KEEP\_K closes only about half its remaining gap, and its per-epoch curve is already plateauing by epoch 8–9 (full curves in `distillation_vision_cifar{10,100}_10ep.csv`), indicating a genuinely harder, slower-converging recovery rather than simple under-training at 5 epochs.

(non-random)  $K$ -derived representation confers no detectable advantage over uninformative noise to the frozen downstream network, even though Section 4.2’s probe shows that representation is just as linearly content-rich as  $V$ ’s.

## 4.2 Testing the addressing/content mechanism directly

Corollary 1 rests on an assumption we had not, until now, tested directly: that the teacher’s  $K$ -projection specialized for *addressing* at some cost to *content*, while  $V$  specialized the other way. We test both halves on the five synthetic-task teachers, where the correct routing target at each position is known exactly (COPY/SUB: self; REVERSE: position  $T-1-i$ ; SWAP: swap the two halves; SORT: the source position is  $\text{arg sort}(x)$ , which is per-sample and data-dependent).

**Addressing.** For each layer, we take the model’s actual attention weights  $a_{ij}$  (re-derived from the true layer-normed block input via a forward hook, not recomputed independently) and check whether  $\text{arg max}_j a_{ij}$  matches the task’s ground-truth source position. On the two genuinely one-to-one permutation tasks this is strong (REVERSE  $1.00 \rightarrow 0.83$  from layer 0 to 1; SWAP  $1.00 \rightarrow 0.78$ ). On SUB and COPY it is near chance ( $\approx 0.02$ – $0.05$ ) at both layers, consistent with Section 3: attention apparently does not implement self-addressing for these, because it does not need to. SORT is *also* near chance under this diagnostic, which we do not read as evidence against SORT needing routing – it collapses under surgery just like REVERSE/SWAP (Table 1) – but as a sign that “does  $\text{arg max}$  match one ground-truth index” is the wrong operationalization for a task whose natural solution is a distributed rank/comparison computation over many positions rather than a hard one-to-one copy. We report this as a limitation of the diagnostic, not a finding about




Figure 2: Per-epoch distillation accuracy for every (task/dataset, surgery-mode) configuration in Tables 1–2 (dashed line: the from-scratch  $Q \neq K=V$  tied baseline). CIFAR-10/CIFAR-100 panels show the full 10-epoch runs; all other panels are the 5-epoch runs (not rerun longer since already converged).

| Task/dataset | KEEP_K | random $c_{kv}$ (mean $\pm$ std, 5 seeds) | $z$    |
|--------------|--------|-------------------------------------------|--------|
| REVERSE      | 0.100  | $0.100 \pm 0.0004$                        | +1.25  |
| SORT         | 0.583  | $0.584 \pm 0.0099$                        | −0.12  |
| SWAP         | 0.088  | $0.100 \pm 0.0010$                        | −12.30 |
| MNIST        | 0.114  | $0.091 \pm 0.0124$                        | +1.92  |
| FMNIST       | 0.084  | $0.118 \pm 0.0278$                        | −1.23  |
| CIFAR-10     | 0.076  | $0.112 \pm 0.0210$                        | −1.72  |
| CIFAR-100    | 0.007  | $0.011 \pm 0.0024$                        | −1.88  |

Table 3: KEEP\_K vs. the random- $c_{kv}$  baseline, with  $z = (\text{keep\_k} - \text{mean})/\text{std}$ . Only SWAP is a large, unambiguous effect; the rest are within about 2 standard deviations of the random baseline’s own spread at  $n = 5$  seeds, i.e. not reliably distinguishable from noise in either direction. SUB and COPY are omitted: every mode, including all 5 random seeds, scores exactly 1.000.

SORT’s mechanism.

**Content.** We fit a closed-form ridge-regression linear probe from  $k_j$  (resp.  $v_j$ ) to  $x_j$ ’s one-hot identity, on a held-out train/test split, and report test-set classification accuracy. The result complicates Corollary 1 as stated: at this model’s width ( $d=256$ ), *both*  $k_j$  and  $v_j$  probe at or near 1.000 on every task and both layers, with no consistent gap (the one partial exception, SORT layer 1, is 1.000 for  $K$  vs. 0.991 for  $V$  – if anything the wrong direction for the hypothesis). So the teacher’s  $K$ -projection has not, in any sense a linear probe can detect, sacrificed the ability to reconstruct token identity. Corollary 1’s premise – that  $K$  is under no pressure to stay content-decodable – does not hold as a loss of *linearly available information*, at least not at this scale.

Read together with the alpha-sweep/random-baseline result above, this points to a more specific mechanism than “ $K$  lost the content”: KEEP\_K confers no detectable advantage over an untrained random baseline on most non-pointwise tasks, and is dramatically worse on one (SWAP; Table 3), even though  $k_j$  is linearly just as decodable as  $v_j$ .  $K$  does not reliably preserve useful downstream functionality despite preserving content, which is consistent with a **basis mismatch** rather than




Figure 3: Zero-shot accuracy as  $c_{kv}$  is swept from KEEP\_V ( $\alpha=0$ ) to KEEP\_K ( $\alpha=1$ ), vs. a random-initialization  $c_{kv}$  baseline (shaded: mean  $\pm$  1 std. over 5 seeds). No retraining anywhere in this figure.

an information loss – though we have not measured downstream calibration to  $K$  vs.  $V$  directly, so we hold this as the best-supported reading rather than an established mechanism:  $k_j$  carries  $x_j$ 's identity in a representation shaped by the addressing objective ( $q_i \cdot k_j$  ranking correctly), not in the representation  $c_{proj}$ , the MLP, and the head were calibrated against – and handing those frozen downstream weights an unfamiliar (if information-rich) basis is, on the evidence in Table 3, no more useful to them than noise. We did not anticipate this when first writing Corollary 1, and we think it is a more precise and more interesting claim than the one we originally set out to test – and it is also consistent with why distillation recovers KEEP\_K slowest (Section 5): re-calibrating a frozen network to a new, unfamiliar basis is a harder fitting problem than recovering from noise-level information loss would be.

### 4.3 K–V representational similarity as a diagnostic for collapse severity

Corollary 1 predicts that the severity of collapse should track how much the teacher's independently-trained  $K$  and  $V$  projections actually diverged. We measure this directly: for each teacher checkpoint, we capture  $c_k(x)$  and  $c_v(x)$  activations on held-out data via forward hooks and compute linear CKA [2] between the two activation matrices, averaged over layers. High CKA means  $K$  and  $V$  computed near-redundant functions (safe to tie); low CKA means they specialized (unsafe to tie).

Across all 9 tasks, Figure 4 shows the predicted negative relationship: SUB and COPY sit at the high-CKA (0.93–0.94) / zero-collapse corner, while REVERSE, SORT, SWAP, and every vision dataset cluster at moderate CKA (0.53–0.75) with real collapse. This global correlation is driven mainly by that between-group contrast – pointwise tasks let  $K$  and  $V$  converge to near-identical functions with no cost, while every routing-dependent task or dataset needs them to specialize at least somewhat.

**How much should  $n = 9$  and that between-group contrast worry us?** A fair amount, and we report the numbers that say so rather than lead with the most favorable one. Pearson's  $r = -0.84$  ( $p = 0.005$ ) has a bootstrap 95% confidence interval of  $[-0.98, -0.18]$  ( $10^4$  resamples) – consistent with a strong relationship, but not tightly enough pinned down at  $n = 9$  to rule out a




Figure 4: Activation-level linear CKA between a teacher’s trained  $K$  and  $V$  projections vs. the accuracy drop under  $K=V$  surgery (averaged over  $\text{KEEP\_K}$  and  $\text{KEEP\_V}$ ), across all 5 synthetic tasks and 4 vision datasets. Pearson  $r = -0.84$  ( $n = 9$ ); see text for why we do not treat this as a validated general predictor.

much weaker one. Spearman’s rank correlation, which does not reward the linear separation between the pointwise cluster and everything else as heavily as Pearson does, is  $\rho = -0.56$  ( $p = 0.12$ , not significant). And within each task-type group alone – synthetic-only ( $n = 5$ :  $r = -0.82$ ,  $p = 0.09$ ) or vision-only ( $n = 4$ :  $r = -0.85$ ,  $p = 0.15$ ) – the correlation is directionally consistent with the global trend but does not reach significance at these small within-group sample sizes. We therefore do not treat CKA as a validated general predictor of collapse severity; it is a useful, cheap diagnostic that lines up with the theory’s qualitative prediction across task types in this dataset, and the honest summary of the evidence for it is weaker than the headline  $r = -0.84$  makes it look.

What the correlation does *not* explain is the ordering *within* the vision group. A natural hypothesis is that harder classification tasks force  $K$  and  $V$  to specialize *more* (lower CKA), which would explain why collapse gets worse from MNIST to CIFAR-100. The data says the opposite: activation CKA *rises* monotonically with dataset difficulty (MNIST 0.53  $\rightarrow$  FMNIST 0.67  $\rightarrow$  CIFAR-10 0.71  $\rightarrow$  CIFAR-100 0.75), even as zero-shot collapse also gets worse. So whatever drives increasing collapse severity across the vision datasets, it is not increasing  $K$ – $V$  specialization – CKA is doing the opposite of what that story requires. We do not have an account for this and are not aware of one; we report it as an open question rather than force a narrative onto it. One plausible (but unverified) reading is that harder classification tasks lean more on the class token aggregating broad, low-precision context from many patches – a use of attention that Proposition 1 restricts regardless of how similar  $K$  and  $V$  are – so collapse severity here may be tracking how much content-routing the task depends on, not how specialized  $K$  and  $V$  happened to become. We leave testing that directly to future work.

**Where in the network it happens.** Every model here has exactly two attention layers, which lets us check layer-by-layer rather than only report the average. The two domains show opposite trends. In the three routing-dependent synthetic tasks, CKA is *lower* in layer 0 than layer

1 (REVERSE 0.54  $\rightarrow$  0.82; SORT 0.56  $\rightarrow$  0.74; SWAP 0.54  $\rightarrow$  0.83):  $K$  and  $V$  specialize early and converge back toward redundancy by the final layer. In vision the trend reverses – CKA is high in layer 0 and drops in layer 1, the layer immediately feeding the CLS readout (MNIST 0.74  $\rightarrow$  0.33; FMNIST 0.77  $\rightarrow$  0.57; CIFAR-10 0.83  $\rightarrow$  0.58; CIFAR-100 0.79  $\rightarrow$  0.70). Layer 0’s CKA across the four vision datasets is not monotonic in difficulty, but layer 1’s is ( $0.33 < 0.57 < 0.58 < 0.70$ , i.e. MNIST < FMNIST  $\lesssim$  CIFAR-10 < CIFAR-100): the across-dataset trend already visible in Figure 4 is carried almost entirely by the layer adjacent to the classification readout. This sharpens, without resolving, the open question above – whatever pushes harder vision datasets toward higher  $K$ – $V$  redundancy concentrates at exactly the layer the class-token-aggregation account above would implicate, which is consistent with that account without being evidence for it over any other.

**Robustness to the similarity measure.** All CKA numbers above use the linear kernel. As a check that the pattern is not an artifact of that choice, we recomputed every measurement with RBF-kernel CKA [2] instead (median pairwise-distance bandwidth; activations subsampled to 2000 tokens per layer for tractability). The global correlation with collapse severity is essentially unchanged ( $r = -0.87$  vs.  $r = -0.84$  for linear CKA,  $n = 9$ , both  $p < 0.01$ ), and the qualitative pattern is preserved throughout, including the vision layer-1 ordering above. We have not evaluated CCA- or mutual-information-based similarity measures; we have no reason to expect a different qualitative story from them, but we have not checked, and leave it to future work rather than claim it.

Weight-space cosine similarity between the raw  $(W_k, b_k)$  and  $(W_v, b_v)$  parameters, by contrast, is uninformative in every task (all values within  $\pm 0.005$  of zero) – as expected, since two linear maps can compute highly redundant *functions* while living in unrelated parameter bases; only the representation-level (CKA) comparison is basis-invariant and captures functional overlap.

## 5 Discussion

Recovery is fast relative to training from scratch: most (task, surgery-mode) configurations reach the from-scratch tied baseline within 1–5 epochs of distillation against the QKV teacher, versus the 10–50 epochs used to train the corresponding checkpoints from scratch, and KEEP\_V/AVG on CIFAR are within a point of that baseline by 10 epochs (Table 2). One fact about the surgery is certain and worth stating plainly: it leaves every parameter *except*  $c_{kv}$  already correctly tuned (Section 1), so distillation is searching for a single linear map rather than re-deriving the whole network. Whether that alone accounts for the observed speed of recovery – as opposed to, say, the KL term against the teacher’s soft logits doing more work than the search-space argument suggests – we have not verified, and we do not offer a convergence-rate account here. KEEP\_K is the clear counter-case to any simple “small search space, therefore fast” story: it starts from the same reduced parameter count as KEEP\_V and AVG but is consistently the slowest and least complete to recover (Table 2). Section 4.2’s basis-mismatch account is more consistent with this than Corollary 1 as originally stated: distillation is not restoring content that KEEP\_K discarded (a linear probe shows the content was never gone), it is re-calibrating  $c_{\text{proj}}$ /MLP/head to a representation in an unfamiliar basis – a harder fitting problem than the search-space argument alone would suggest.

## 6 Related work

**Attention as constrained computation.** Tsai et al. [4] reframe self-attention generally as a kernel smoother:  $y_i$  is a weighted average of value vectors, with the softmax scores acting as a data-dependent kernel. That framing already implies Proposition 1 as a special case once the

vectors being smoothed and the vectors used to produce the kernel weights are forced to be the same object ( $k = v$ ): the kernel-smoother view says the output lies in the convex hull of whatever is being smoothed, and our contribution is narrower and architecture-specific – showing what that hull becomes, and why it is a genuine bottleneck for some tasks and not others, exactly when  $K = V$ .

Two other lines of work probe different senses in which attention is “limited,” and it is worth being precise about how ours differs from each. Dong et al. [5] show that stacked self-attention without skip connections or MLPs collapses token representations to rank 1 doubly exponentially with depth – a degeneracy that compounds across many layers of attention interacting with each other. Proposition 1’s obstruction is unrelated to depth: it is a single-layer statement about the span reachable by one attention block once  $K = V$ , holds regardless of how many layers are stacked, and does not require removing the skip connections or MLPs that Dong et al. show arrest their collapse (our checkpoints keep both throughout). Raganato et al. [6] show that replacing most encoder attention heads with fixed, non-learned, position-based patterns barely hurts translation quality – i.e. the specific values of the attention weights  $a_{ij}$  are often not doing much specialized work. That is an empirical, training-time finding about how much freedom the *weights* need; ours is an architectural, weight-independent statement about what the *value space* can carry once  $K = V$  – Proposition 1 holds for every possible  $a_{ij}$ , learned or fixed, so the limitation we describe is structural rather than something a better-trained or differently-patterned attention could route around.

**Representational similarity.** Our CKA measurements follow Kornblith et al. [2], and the finding that raw weight-cosine similarity is uninformative while activation-level CKA is not (Section 4) is the expected consequence of CKA’s basis invariance – precisely the property it was designed to have.

**Distillation.** Our recovery protocol (hard-label cross-entropy plus KL divergence against soft teacher logits,  $T = 2$ ) follows Hinton et al. [3]. That paper, and most of the compression literature built on it, uses distillation to shrink a model while preserving its architecture class; here distillation instead repairs a capacity loss introduced by an architectural edit made *after* training (Section 4) – a different use of the same mechanism.

**KV-sharing architectures.**  $K=V$  tying is a special case of a broader, well-established family: multi-query attention shares  $K/V$  across all heads [7], and grouped-query attention shares them across groups of heads [8]; the base paper’s own sweep includes exactly these variants (GQA/MQA and their  $Q \neq K=V$  combinations) trained jointly from scratch at 300M and 1.2B scale. None of that is new, and we do not claim otherwise. What this note adds is orthogonal to the architecture choice itself: a mechanism for *why* imposing that same sharing structure *after* independent training is a qualitatively different operation from choosing it *before* training starts, and a cheap, testable account (Sections 2, 4.2) of when the difference is severe enough to matter. A reviewer’s fair objection – “ $K/V$ -sharing is not new; this is already standard in efficient attention” – is correct about the architecture and, we think, orthogonal to what we are claiming about the post-hoc/joint-training distinction.

## 7 Conclusion

The result we would ask a reader to take away is not “ $K = V$  is fundamentally deficient” –  $Q \neq K=V$  trained jointly from scratch works fine (the base paper’s headline result), and two of our five synthetic tasks are entirely untroubled by the surgery. The result is that **post-hoc merging of independently-trained representations is a different operation from a joint architectural constraint, with its own, separately characterizable failure mode**. Joint training

lets addressing and content pressures negotiate a shared representation from the start; post-hoc merging imposes that constraint externally on two representations that specialized apart under no such pressure, and Proposition 1 gives a hard, narrow ceiling on what a single merged attention sub-block – not the whole network – can then carry (Remark 2).

Where that ceiling ends up mattering turned out to need more than the mechanism we started with. Corollary 1’s original story –  $K$  addresses,  $V$  carries content, so  $K$  should be the less content-decodable of the two – does not survive a direct probe:  $K$  and  $V$  are equally linearly decodable (Section 4.2). The evidence is more consistent with a more specific, basis-mismatch account than with Corollary 1 as originally stated – though we have not directly measured downstream calibration to  $K$  vs.  $V$ , so we offer this as the best-supported reading, not an established mechanism: measured against a random-projection baseline’s own 5-seed spread rather than a single number, KEEP\_K confers no detectable advantage over noise on most non-pointwise tasks and is dramatically worse on one (Table 3), even though its content is, by the same section’s probe, fully present. K–V activation CKA is a useful and cheap diagnostic that tracks collapse severity across task types in our data, but we do not oversell it as a validated general predictor: the correlation is weaker under rank statistics than under Pearson’s, is not separately significant within either task-type group at our sample sizes, and runs in the opposite direction of naive expectation *within* the vision datasets for reasons we do not have an account for (Section 4) – three separate reasons for a future study with more than  $n = 9$  tasks to revisit it. When the ceiling does matter, a short knowledge-distillation pass against the original model is a comparatively cheap and nearly-always-sufficient fix, with KEEP\_K consistently the slowest to recover – which the basis-mismatch account predicts, since it is the mode with a genuinely unfamiliar representation to re-calibrate against, not merely a smaller one.

## References

- [1] A. Kayyam et al. *Do Transformers Need Three Projections?* ICML (poster), 2025.
- [2] S. Kornblith, M. Norouzi, H. Lee, G. Hinton. *Similarity of Neural Network Representations Revisited*. ICML, 2019.
- [3] G. Hinton, O. Vinyals, J. Dean. *Distilling the Knowledge in a Neural Network*. arXiv:1503.02531, 2015.
- [4] Y.-H. H. Tsai, S. Bai, M. Yamada, L.-P. Morency, R. Salakhutdinov. *Transformer Dissection: An Unified Understanding for Transformer’s Attention via the Lens of Kernel*. EMNLP, 2019.
- [5] Y. Dong, J.-B. Cordonnier, A. Loukas. *Attention Is Not All You Need: Pure Attention Loses Rank Doubly Exponentially with Depth*. ICML, 2021.
- [6] A. Raganato, Y. Scherrer, J. Tiedemann. *Fixed Encoder Self-Attention Patterns in Transformer-Based Machine Translation*. Findings of EMNLP, 2020.
- [7] N. Shazeer. *Fast Transformer Decoding: One Write-Head is All You Need*. arXiv:1911.02150, 2019.
- [8] J. Ainslie, J. Lee-Thorp, M. de Jong, Y. Zemlyanskiy, F. Lebrón, S. Sanghai. *GQA: Training Generalized Multi-Query Transformer Models from Multi-Head Checkpoints*. EMNLP, 2023.